

Jumpshot , Vampir and TAU

Rosa Filgueira

Data Intensive Research, School of
Informatics, University of Edinburgh

rosa.filgueira@ed.ac.uk



Jumpshot: Basic Information

- Name: MPE/Jumpshot
- Developer: Argonne National Laboratory
- Current versions:
 - MPE 1.26
 - Jumpshot-4
- Website:
<http://www-unix.mcs.anl.gov/perfvis>

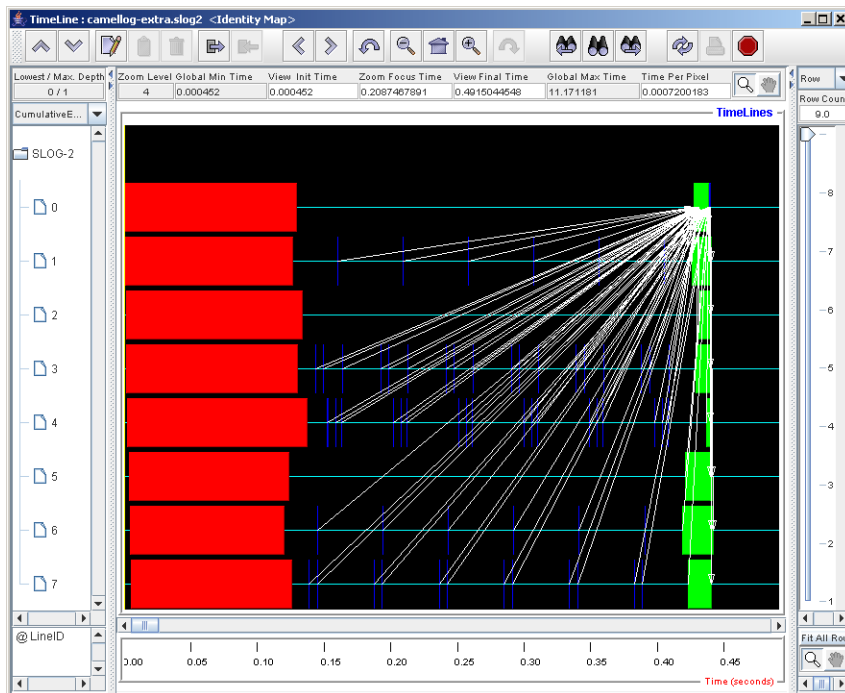
What Is MPE/Jumpshot?

- **MPE** stands for "Multi-Processing Environment". It is a suite of performance analysis tools for MPI programs.
- Creates a log file(s) while the MPI program executes, then provides tools to analyze and view the log files after the execution.
 - **clog2 format** → we need to convert to **slog2 format**
- *Jumpshot* is a Java-based GUI tool for viewing one of these log files:
 - A visualization tool for logfiles (slog2) created by the MPE package
 - Written in Java (crossplatform)
 - Provides a "time line" (GANTT) view of MPI and program events
 - Also has basic search and summary (histogram) functionality

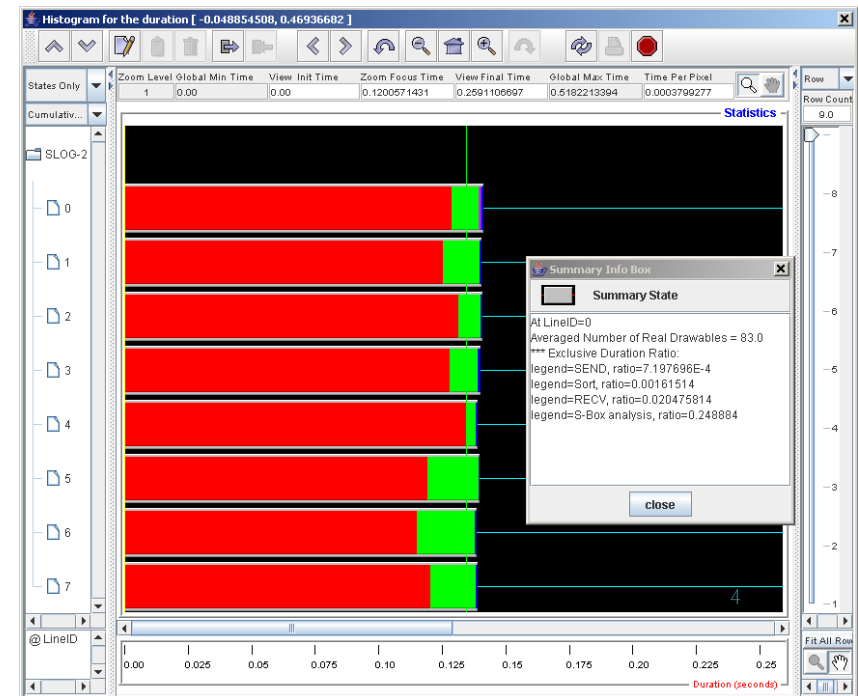
Jumpshot Overview

- Jumpshot-4 supports two types of visualizations for metrics
 - Timeline (left, botto,)
 - Histogram (right, bottom)
- Visualization is dependant on SLOG-2 format and Data model

Timeline view



Histogram view



Hello Word

```
#include "stdio.h"
#include <stdlib.h>
#include <mpi.h>
#include <string.h>
#include <unistd.h>
int main(int argc, char *argv[])
{
    int rank, size, i;
    int type = 0;
    MPI_Status status;
    char message[20];
    char hostname[1024];
    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD,&size);
    MPI_Comm_rank(MPI_COMM_WORLD,&rank);
    hostname[1023] = '\0';
    gethostname(hostname, 1023);
    printf("My Hostname proc %d : %s\n",rank,hostname);
    if (rank == 0) {
        strcpy(message,"Hello, world");
        for (i = 1; i < size; i++) MPI_Send(message,
        13,MPI_CHAR,i,type,MPI_COMM_WORLD);
    }
    else
        MPI_Recv(message,20,MPI_CHAR,0,type,MPI_COMM_WORLD,&status);
    printf("Message from node %d: %13s\n", rank,message);
    MPI_Finalize();
}
```

How to use MPE + Jumpshot

- Compile a MPICH program with MPE Logging Library:
 - `[rosa@moore mpi_exercise]$ mpicc -g -c hello.c`
 - `[rosa@moore mpi_exercise]$ mpicc -o hello hello.o -lmpe -lmpe -lm`
- Run the application:
 - `qsub run.sh`
 - It should generate a log file `hello.clog2`
- Convert log format from CLOG2 to SLOG2:
 - `[rosa@moore mpi_exercise]$ /root/rosa_filgueira/slog2sdk-1.2.6/bin/clog2TOslog2 hello.clog2`

Convert to SLOG2 format

```
[rosa@moore mpi_exercise]$ /root/rosa_filgueira/slog2sdk-1.2.6/bin/clog2TOslog2  
hello.clog2
```

```
GUI_LIBDIR is set. GUI_LIBDIR = /root/rosa_filgueira/slog2sdk-1.2.6/lib
```

```
SLOG-2 Header:
```

```
version = SLOG 2.0.6
```

```
NumOfChildrenPerNode = 2
```

```
TreeLeafByteSize = 65536
```

```
MaxTreeDepth = 0
```

```
MaxBufferByteSize = 715
```

```
Categories is FBinfo(478 @ 823)
```

```
MethodDefs is FBinfo(0 @ 0)
```

```
LineIDMaps is FBinfo(232 @ 1301)
```

```
TreeRoot is FBinfo(715 @ 108)
```

```
TreeDir is FBinfo(38 @ 1533)
```

```
Annotations is FBinfo(0 @ 0)
```

```
Postamble is FBinfo(0 @ 0)
```

```
Number of Drawables = 21
```

```
Number of Unmatched Events = 0
```

```
Total ByteSize of the logfile = 3040
```

```
timeElapsed between 1 & 2 = 13 msec
```

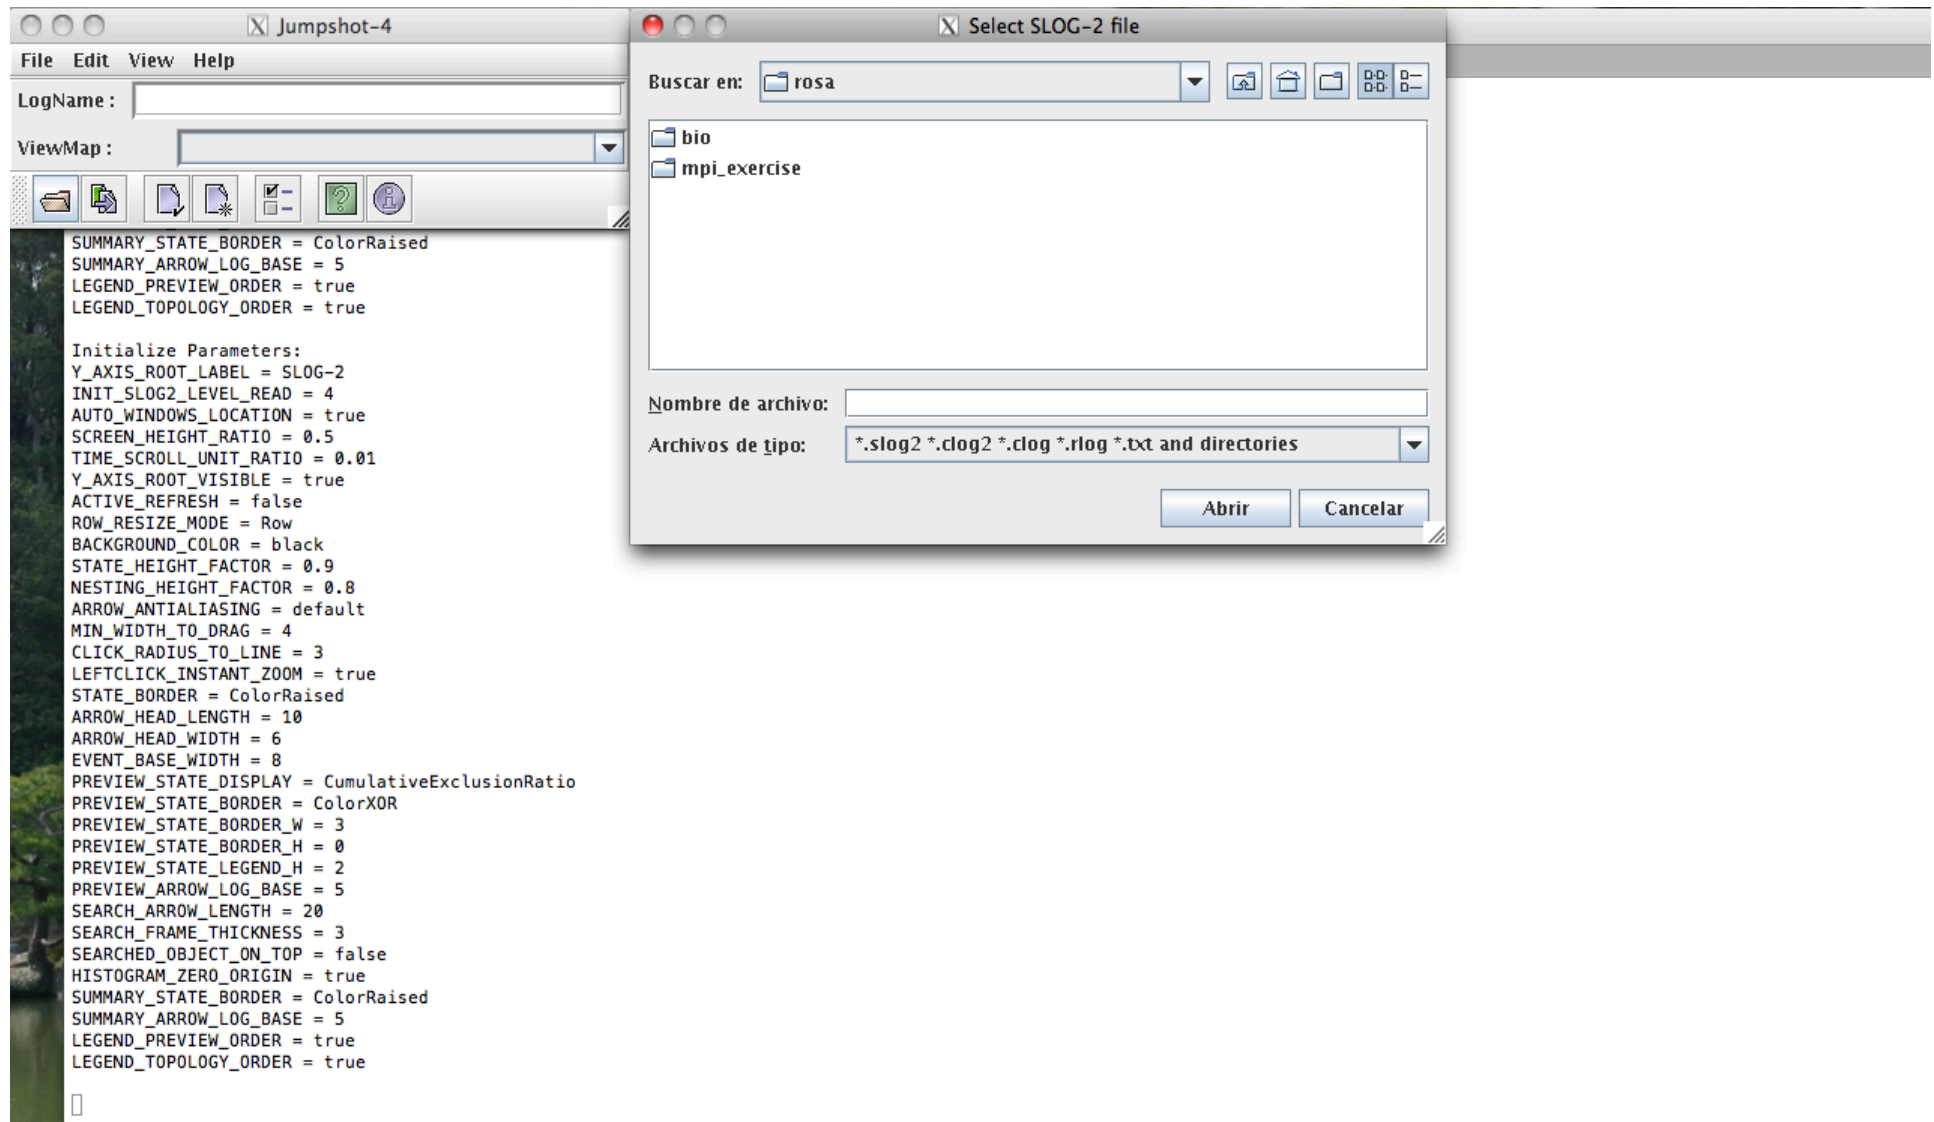
```
timeElapsed between 2 & 3 = 48 msec
```

```
[rosa@moore mpi_exercise]$
```

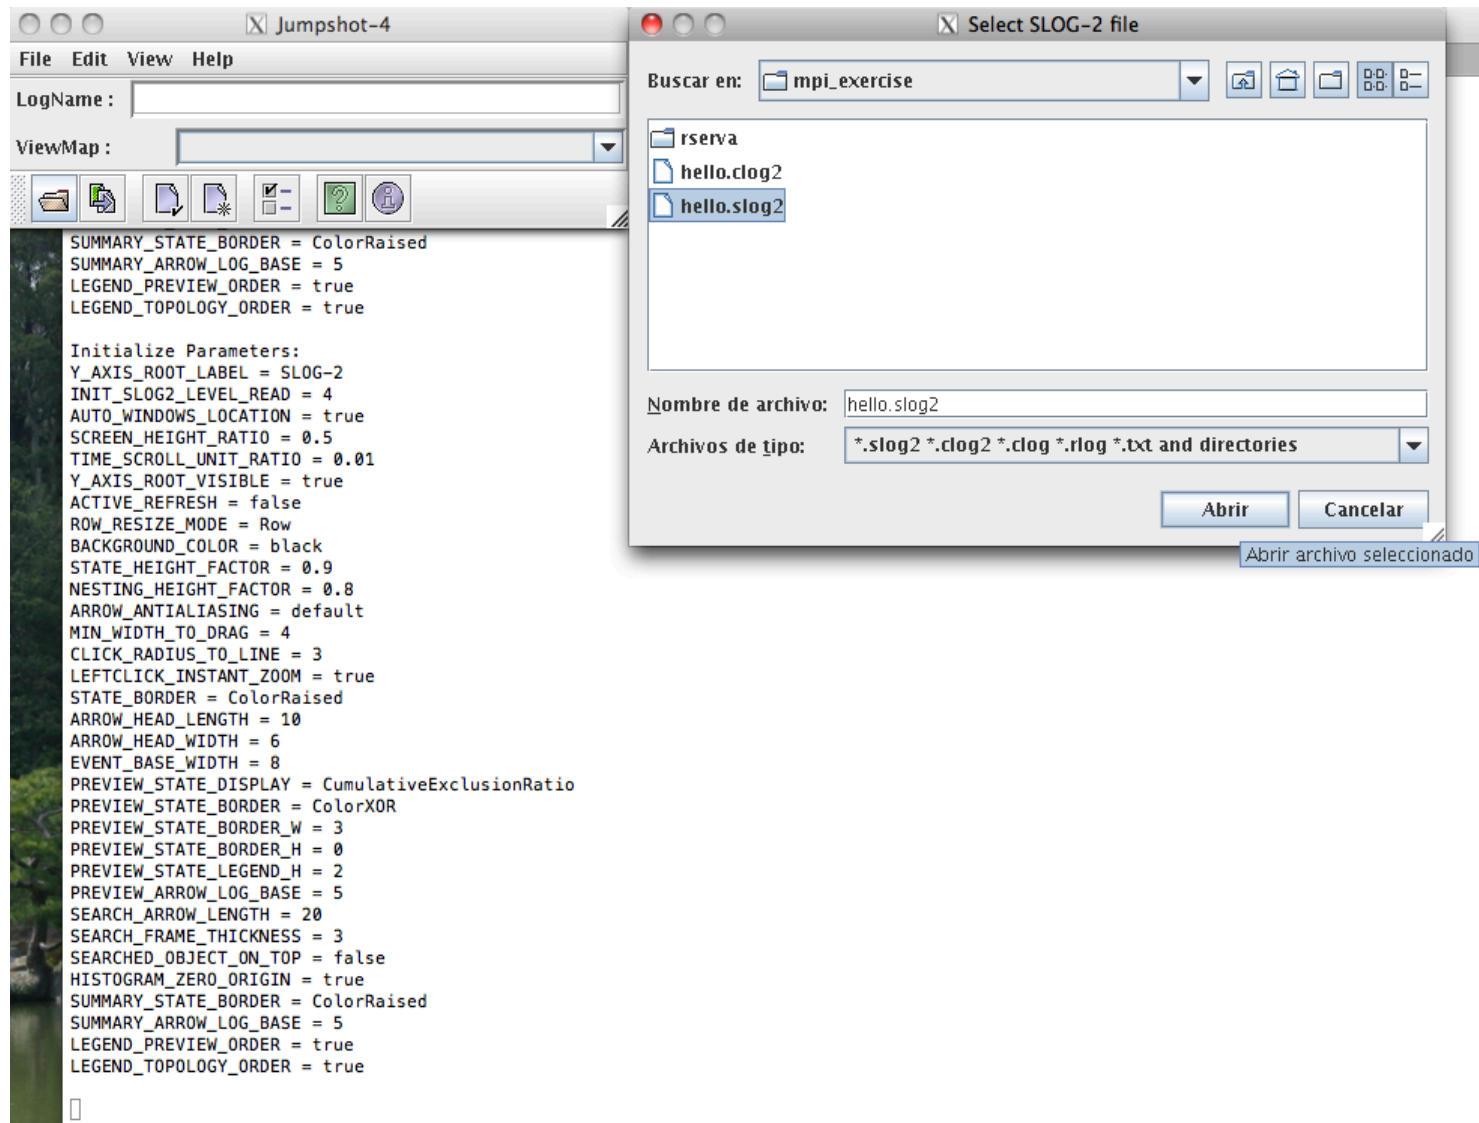
MPE + Jumpshot

- Start Jumpshot-4:
 - New terminal: `ssh -X user@moore.udl.net`
 - **`/share/apps/jumpshot/bin/jumpshot`**
 - First time that the user starts jumpshot, the application generates a configuration file and asks to save it. → Yes.

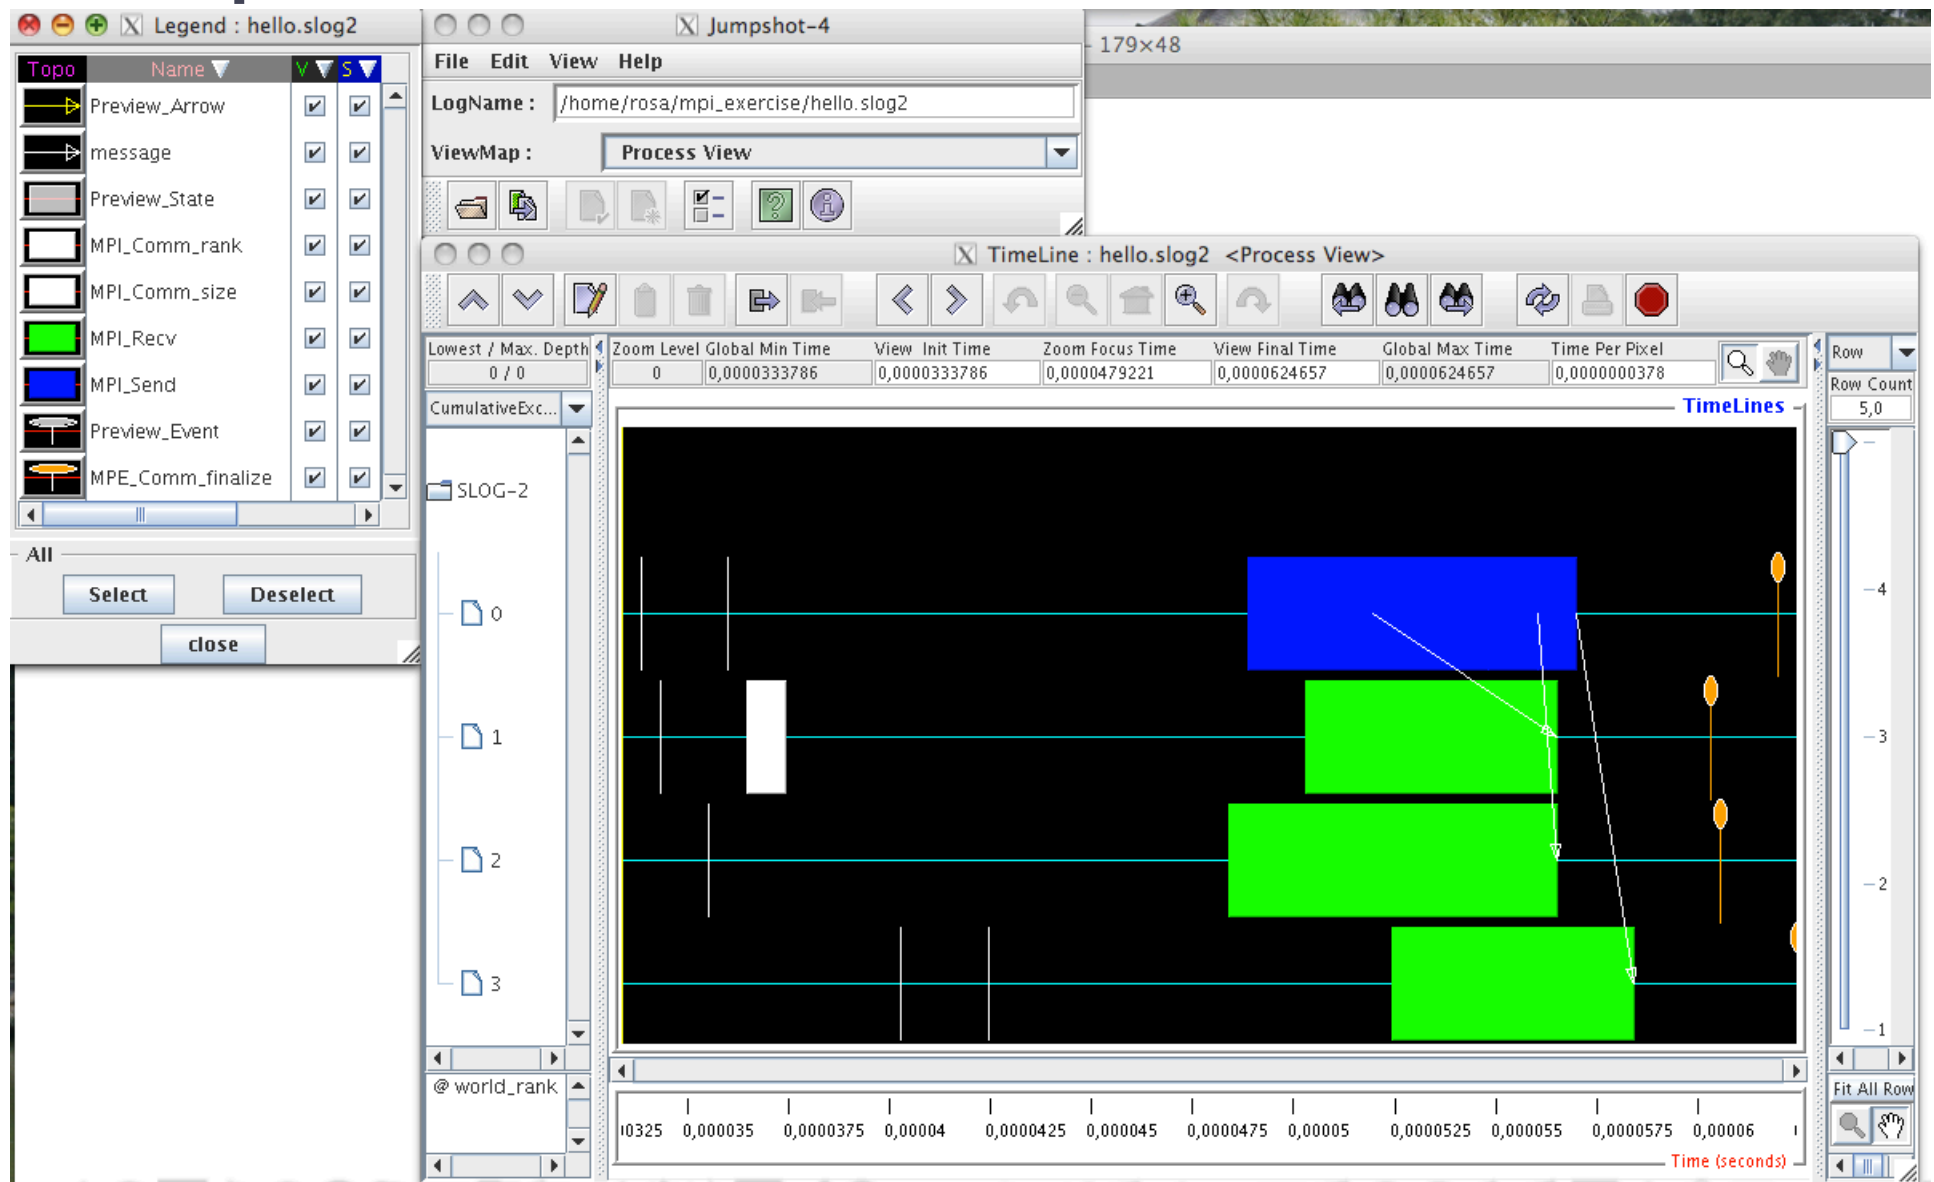
JumpShot-4



JumpShot-4



JumpShot-4



JumpShot-4

Open clog2ToSlog2 Configure preferences

The screenshot shows the JumpShot-4 application window. The interface includes a menu bar (File, Edit, View, Help), a toolbar with icons for file operations and viewing, and a main display area showing a timeline of events. The left sidebar contains a list of event types with checkboxes for selection. The bottom status bar shows the current time in seconds.

Annotations and actions:

- Open clog2ToSlog2:** A red arrow points to the 'Open' icon (a folder) in the toolbar.
- Configure preferences:** A red arrow points to the 'Preferences' icon (a gear) in the toolbar.
- Zoom in-out:** A blue arrow points to the 'Zoom In' and 'Zoom Out' icons in the toolbar.
- Close Window:** A blue arrow points to the 'Close' icon (a red X) in the toolbar.
- Edit Events:** A blue arrow points to the 'close' button in the 'All' section of the left sidebar.

Edit Events:

- Colors
- Disable or enable event

Timeline Data:

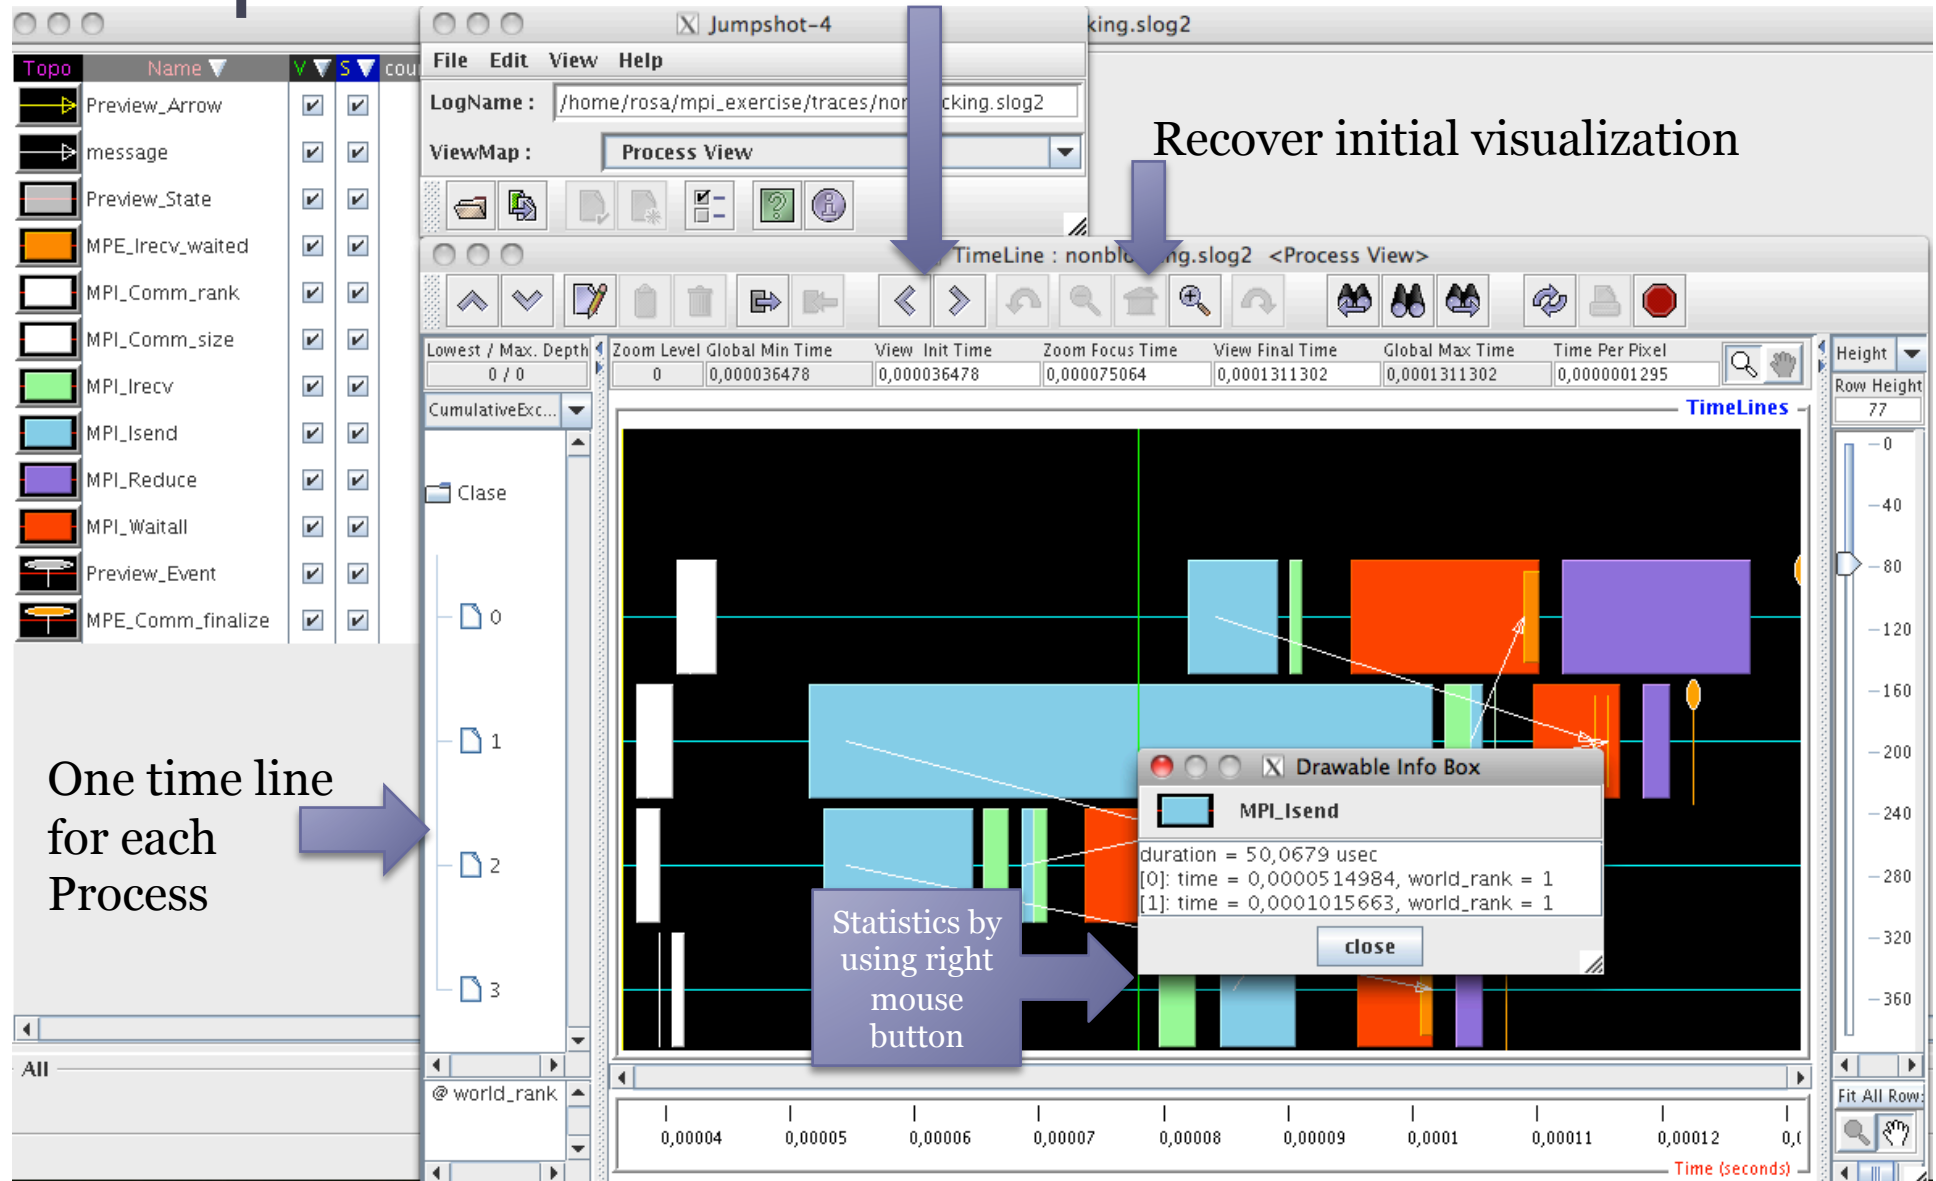
Lowest / Max. Depth	Zoom Level	Global Min Time	View Init Time	Zoom Focus Time	View Final Time	Global Max Time	Time Per Pixel
0 / 0	0	0,0000333786	0,0000333786	0,0000479221	0,0000624657	0,0000624657	0,0000000378

Timeline Labels: SLOG-2, TimeLines, Time (seconds)

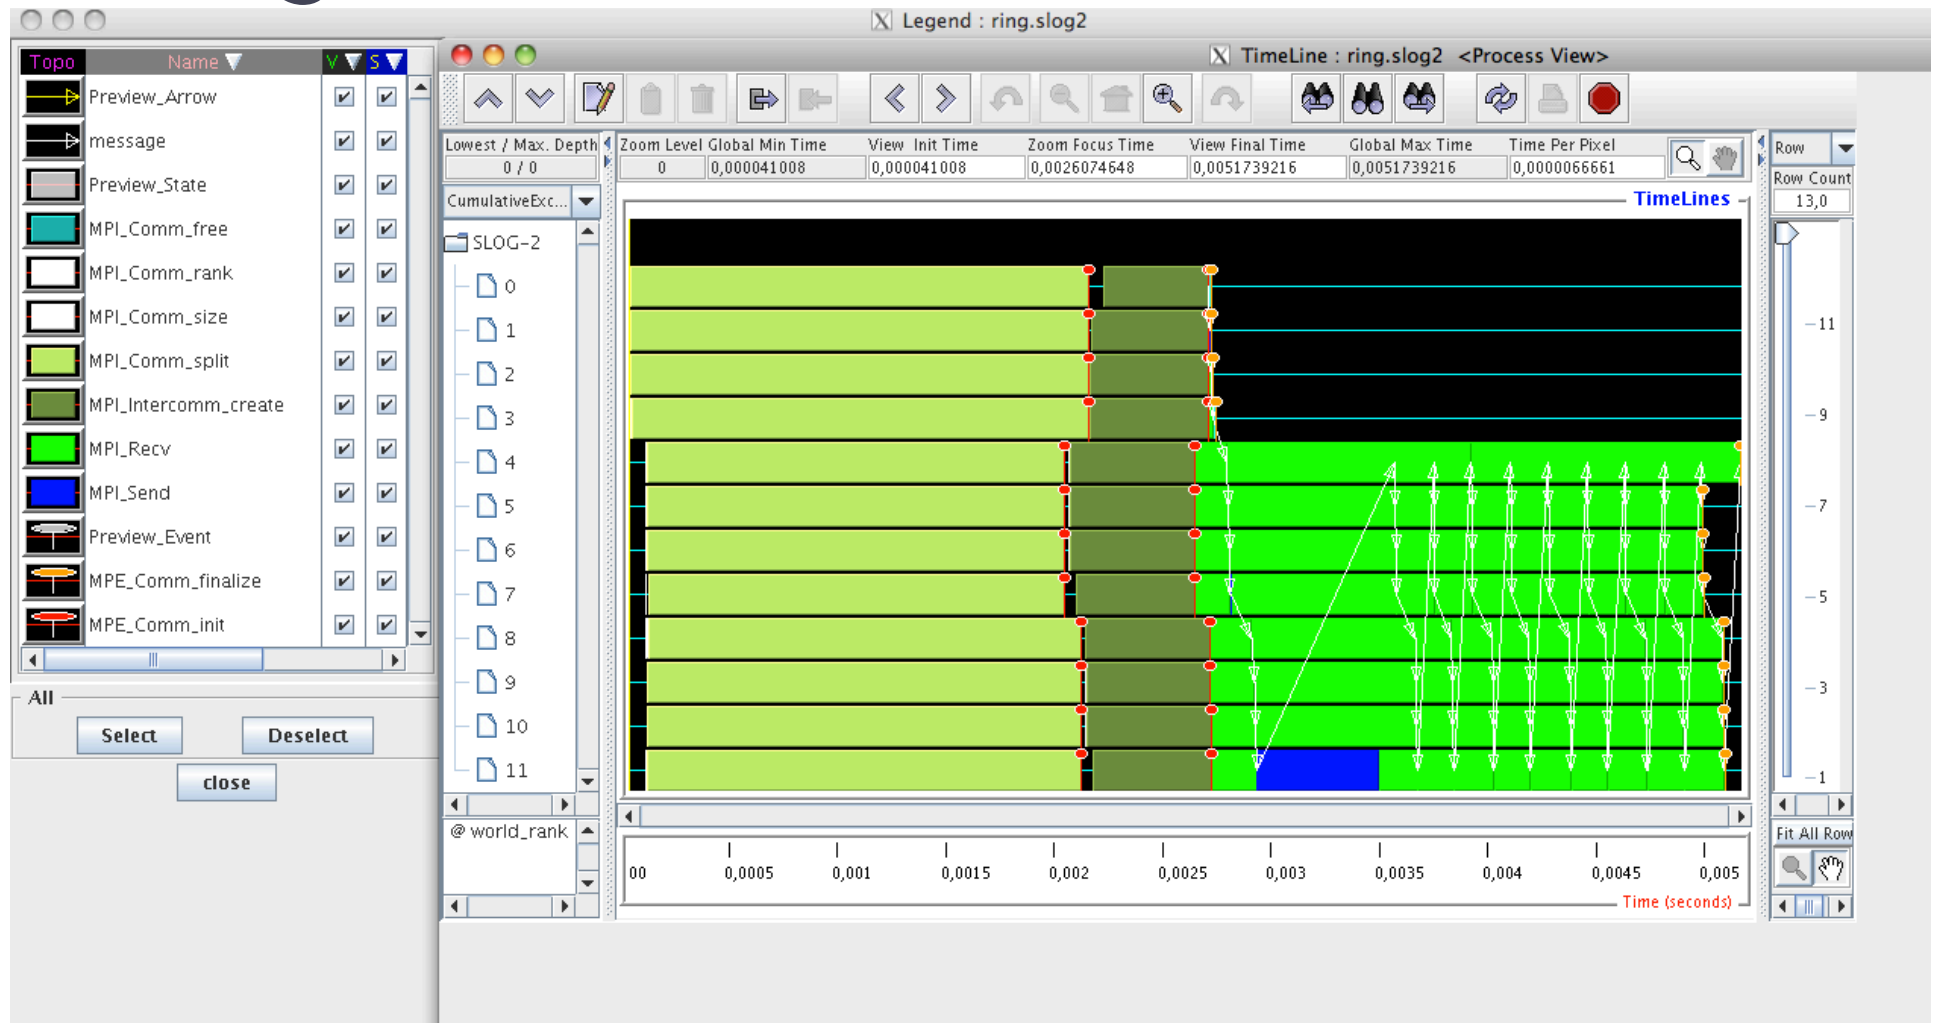
JumpShot-4

Move through the time line

Recover initial visualization

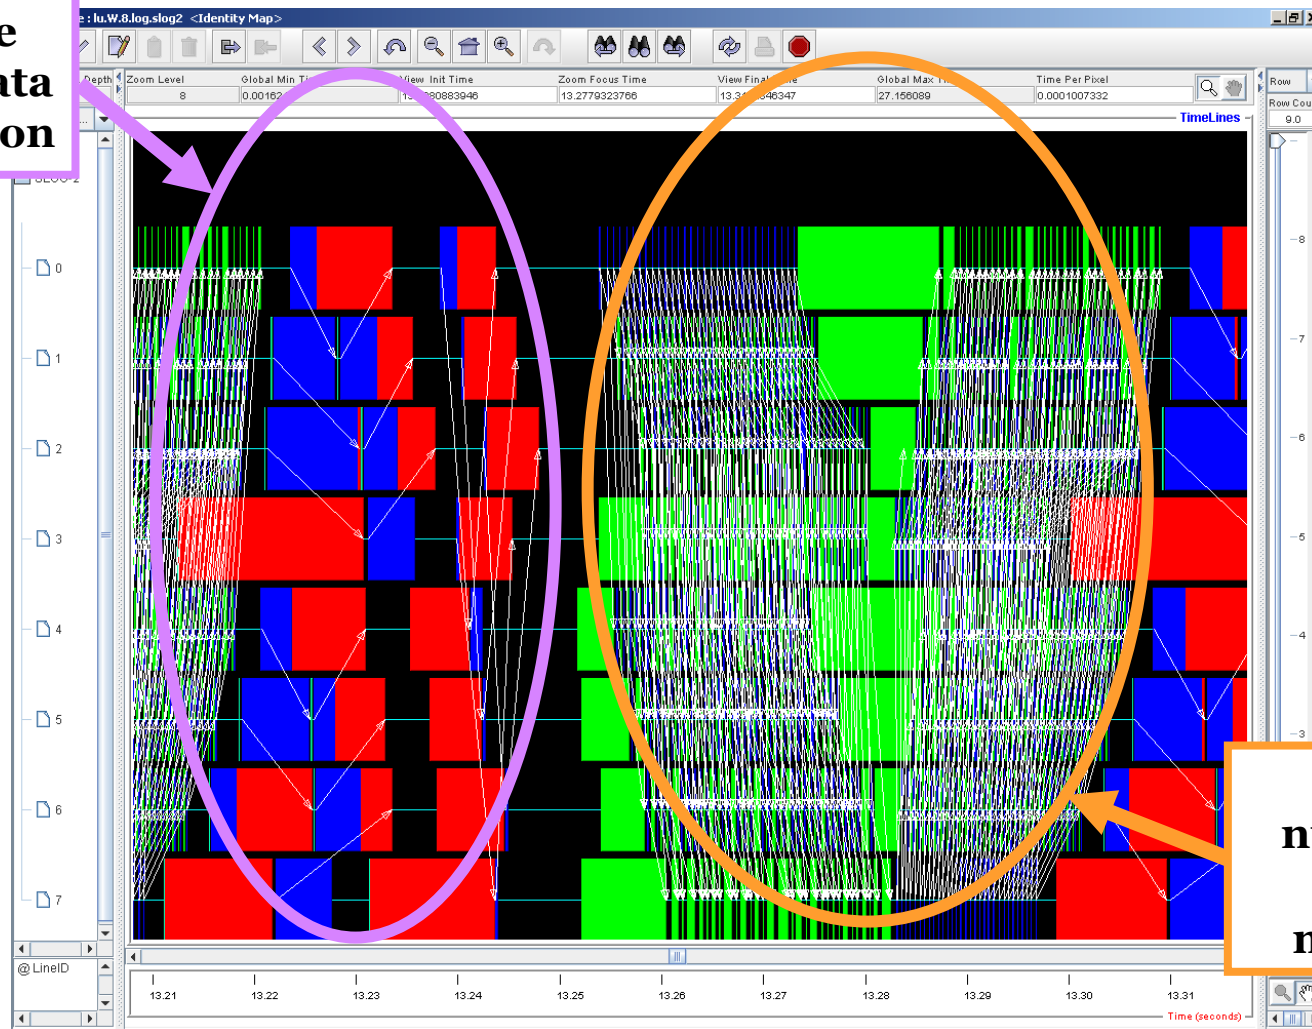


Ring.c



JumpShot-4

**Much time
taken for data
redistribution**



**Large
number of
small
messages**



Vampir

- Performance Analysis Framework
 - Easy to use performance analysis framework for parallel programs
 - Graphical data representation enables detailed understanding of dynamic processes on massively parallel systems
 - In-depth event based analysis of parallel run-time behavior and interprocess communication
 - Identification of performance problems and bottlenecks
 - Linux-based PCs and Clusters, SGI, IBM, SUN, NEC, HP, Apple, ..



Vampir

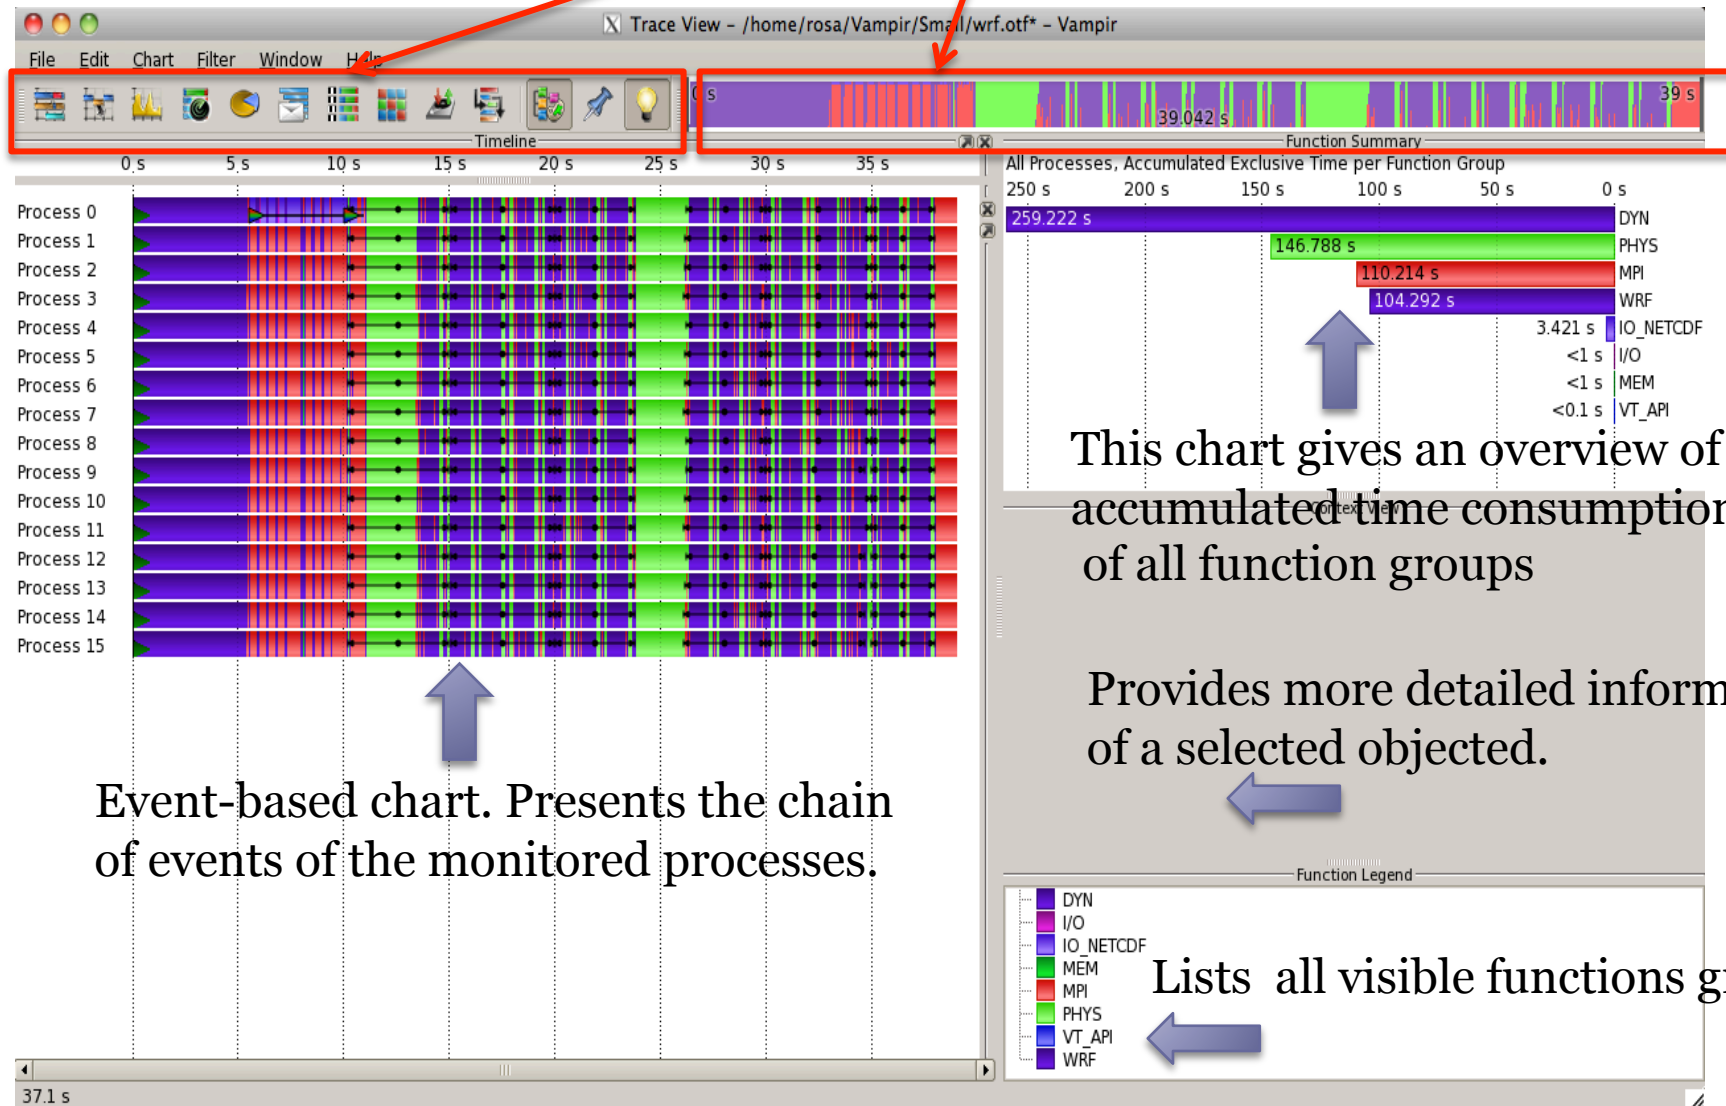
- **Vampir Features**

- Powerful zooming and scrolling in all displays
- Adaptive statistics for user selected time ranges
- Filtering of processes, functions, messages, collective operations
- Hierarchical grouping of threads, processes, and nodes
- Support of source code locations
- Integrated snapshot and printing for publishing
- Customizable displays

Vampir

Tool bar

Zoom Tool bar





TAU

- TAU Performance Systems is a portable profiling and tracing toolkit for performance analysis of parallel programs written in Fortran, C, C++, Java, and Python.
- TAU (Tuning and Analysis Utilities) is capable of gathering performance information through instrumentation of functions, methods, basic blocks, and statements.
- TAU's profile visualization tool, paraprof, provides graphical displays of the performance analysis results, to help the user visualize the collected data.



TAU instrumentation

- **Dynamic Instrumentation:** No change to your application's source or binary code is necessary. Using the `tau_exec` command before your executable generates both profile and trace data.
- **Compiler-Based Instrumentation:** TAU can also be used to compile your application using wrapper scripts for `mpicc`, `mpicxx`, and `mpif90`. Tau instruments the binary code without changing the source. Using the `tau_exec` command before the TAU-compiled executable generates profile and trace data. Data is collected on both MPI calls and your application's function calls.



Visualize Information

- Jumpshot: To visualize traces
- Paraprof: To visualize profiles
 - Tool provided by TAU
 - GUI (Java)
 - Text and graphics
 - Options
 - `--pack <file>`
 - Pack the data in a unique file (no GUI)
 - `paraprof -pack app.ppk`

Dynamic Instrumentation

- Compile
 - `mpicc -o hello hello.c`
- Dynamic instrumentation
 - Modify the “run.sh” script to insert
 - `export TAU_TRACE=1` # This line is needed only for tracing.
 - `export TAU_PROFILE=1` # This line is needed only for profiling.
 - `mpiexec -f $MPICH_MACHINES -n $NSLOTS tau_exec -T PROFILE ./hello`



Generating traces and profiles

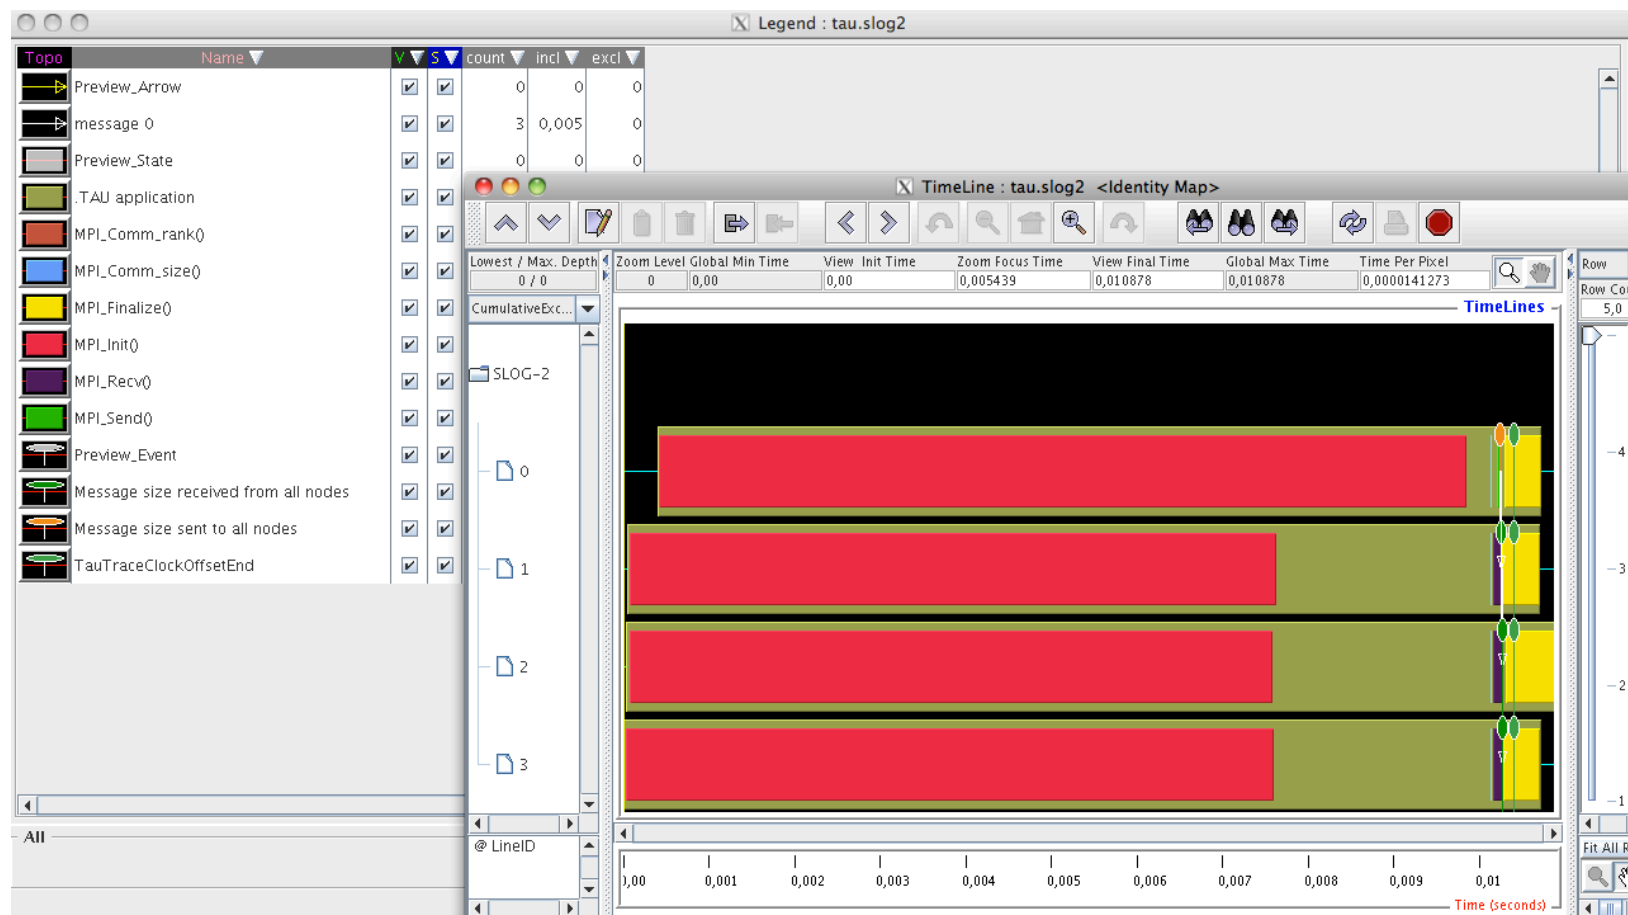
- Run the application
 - `qsub run.sh`
 - The script will generate two types of data:
 - Profile data:
 - `profile.node.o.o` (One per node)
 - Trace data:
 - `tautrace.node.trc` (One per node)
 - `events.node.edf` (One per node)



Visualization of traces

- We need to pack all the tautraces.* in one file, and all the events.* in other file:
 - `tau_treemerge.pl`
 - Generate a unique tautrace file and unique events file
- Generating slog2 file:
 - `tau2slog2 tau.trc tau.edf -o tau.slog2`
- Finally, the tau.slog2 can be visualized by using Jumpshot application

Visualization of traces






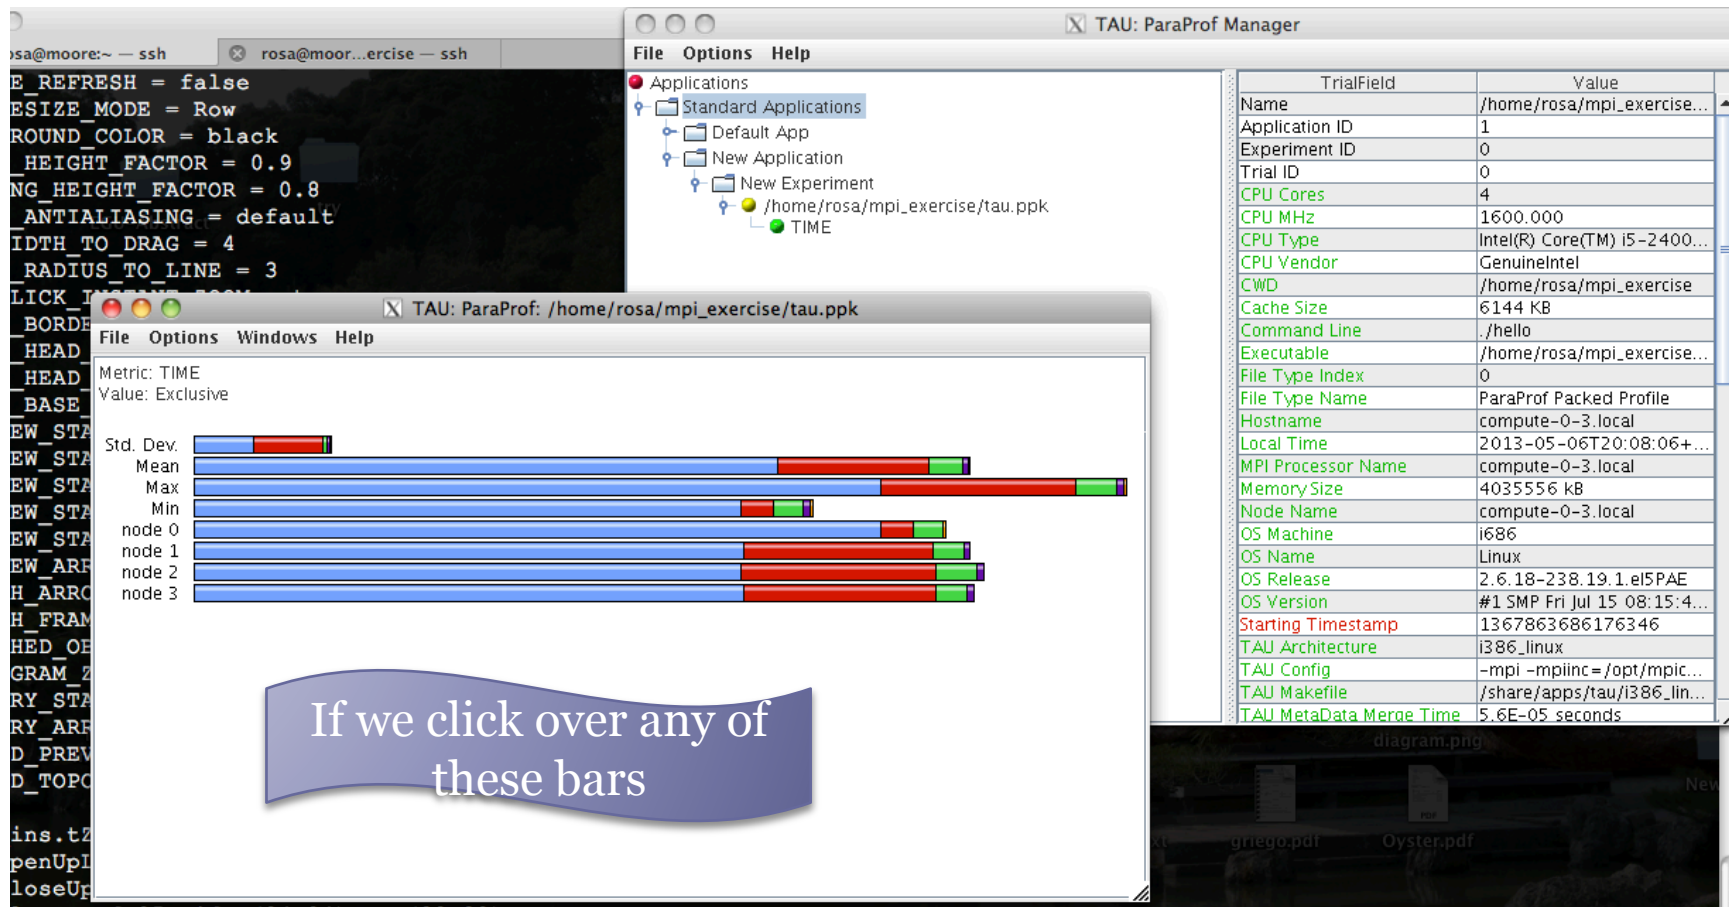
Visualization of profiles

- We need to pack all the profiles.* files in a unique one:
 - `paraprof --pack tau.ppk`
- Now, we can visualize the tau.ppk file by using Paraprof
 - Open a new terminal by using `-X`
 - `/share/apps/tau/i386_linux/bin/paraprof`

Visualization of profiles



Visualization of profiles



Visualization profiles

